Motif-Aware Graph Embeddings

Hoang Nguyen and **Tsuyoshi Murata**
School of Computing, Computer Science Department
Tokyo Institute of Technology

Abstract

In this paper, we propose our motif-aware approaches to the unsupervised network embedding and semi-supervised network labeling task. Our first algorithm is an unsupervised network embedding algorithm which uses the most statistically significant network motif as the guiding pattern for random walks to generate network context. We then use a Skipgram neural network to learn the latent network node representations from the generated context via Noise Contrastive Estimation. The second algorithm employs the Graph Convolution Network model on motif Laplacian matrices to inject the higher-order network structure into the neural network. Both of our algorithms utilize the higher-order organization (i.e. motifs organization) of complex networks. We demonstrate the effectiveness of our algorithms in comparison with other state-of-the-art network embedding algorithms.

1 Introduction

1.1 Complex network and machine learning

Network modelings have been an essential tool for a wide range of scientific fields Newman [2010]; Bader *et al.* [2003]; Milo *et al.* [2002]; Benson *et al.* [2016]. Based on the system's network structure, scientists can make predictions and explanation about the system's behavior. For example, in biology, the study of neuronal system connectivity indicated that the arrangement of neurons is optimized for short processing paths rather than wiring lengths Kaiser and Hilgetag [2006]. Similarly, social network analysis provides community structures well as social interaction patterns West *et al.* [2014]; Barabási [2014]. However, along with the information explosion, analyzing large network-structured datasets poses a great challenge for traditional network analysis methods in term of scalability and complexity. To deal with such challenge, one promising approach is to apply machine learning methods (especially deep learning) to network problems.

Figure 1: Learning a latent representation of nodes from the network structure. In this figure, the example network is embedded to a 3-dimensional real vector space.

Bridging the gap between network-structured data and typical data structure for machine learning is also a challenge. Due to the irregularity in the network-structured data, it is desirable to have a *meaningful* network representation for machine learning applications. Learning network representation in real vector spaces can be viewed as manifold learning (non-linear dimensionality reduction), and commonly known as *network embedding* in the literature. In this report, we propose the use of network motifs into learning a high-quality network embedding. The *embedding quality* in this context is justified by how well a common machine learning model performs on the learned embeddings.

1.2 Motifs in complex networks

There are three scales of network analysis: macroscopic, mesoscopic, and microscopic. In the macroscopic scale, we consider a network as a whole to study its macro-properties such as robustness Callaway *et al.* [2000], or dynamics Barabási [2014]. In contrast, the microscopic scale studies the pair-wise interactions between nodes in a network which is specific to a given system Newman [2010]. In between macroscopic and microscopic, the mesoscopic scale considers the network is a composition of subgraphs. In many fields of research, especially computational biology, the mesoscopic components are called *motifs*, and it is common to think of them as building blocks of a complex system Milo *et al.* [2002].

Definition 1.1. *Network motif*: Given a graph $G = (V, E)$, define a subgraph $G' = (V', E')$ with $V' \subseteq V$; $E' \subset E$ s.t. $i, j \in V', \forall e_{ij} \in E'$ and $|V'| \ll |V|$. Recurring subgraph is called *network motif* when they are statistically significant.

Figure 2: Size-3 motifs

Figure 3: Size-4 motifs and anchor set example

Also referred as the higher-order organization by Benson *et al.*, network motifs are believed to represent the underlying mechanism of a complex system Alon [2007]; Mangan and Alon [2003]. For instance, the directional bi-fan motif (Figure 3: m4-1) and its simplified unidirectional version (m4u-3) are crucial in a citation network. This bi-fan motif is also intuitively sensible in citation network as it represents the citation mechanism. The correlation of recurring subgraphs and system functionality has been studied extensively in biological systems such as transcription networks Mangan and Alon [2003] and brain networks Van Den Heuvel and Pol [2010]; Honey *et al.* [2007]. As networks motifs have been recognized as the fundamental building block of a complex network, using them as a structural guidance for machine learning on network data is expected to yield positive improvements.

Our main idea in this paper is to construct the motif co-occurrence matrix from a given network, and use it as: 1. An adjacency matrix describing a motif network for random walks; 2. A mean to compute motif Laplacian and Fourier basis for the graph convolution operation. Section 2 describes the related work on network motif conductance and network embedding. We give detail of our algorithms in section 3. The experimental setup and results are given in section 4 and 5 respectively. We discuss the relationship between our two proposed algorithms and their limitations in section 6.

2 Related work

In this section, we introduce the recent developments regarding unsupervised network embeddings, semi-supervised node labeling, and motif analysis.

2.1 Unsupervised Network Embedding

Traditionally, network embedding can be obtained via graph factorization methods. However, matrix factorization methods such as Spectral Clustering or Non-negative Matrix Factorization are shown to be unscalable due to the complexity of the algorithms Perozzi *et al.* [2014]; Belkin and Niyogi [2001]. Recently, several feasible network embedding algorithms have been proposed such as *Deepwalk* Perozzi *et al.* [2014] or *node2vec* Grover and Leskovec [2016]. These network embeddings algorithms learn high-quality node representation while having low time complexity compared to traditional methods. In the context of graph embedding, we justify the *embedding quality* by how well a common machine learning model performs on the learned embeddings.

Based on the Skipgram model Mikolov and Dean [2013] in natural language processing, Perozzi *et al.* proposed their scalable graph embedding algorithm named *Deepwalk*. Their results on node classification proved the effectiveness of *Deepwalk* in learning a lower dimensionality representation of a complex network. Subsequent work to *Deepwalk* further improved node classification accuracy by modifying graph context generation process Tang *et al.* [2015]; Cao *et al.* [2015]; Yang *et al.* [2016]; Grover and Leskovec [2016]. Generally, these algorithms optimize the following objective:

$$\mathcal{O} = \underset{W^{\text{emb}}}{\arg\max} \left(\ell_{\text{random walk}} \right), \qquad (1)$$

where $\ell_{\text{random walk}}$ represents the log-likelihood of the network context generated by random walks on the given network. We have the pairwise potential between a *context* vertex v_c and a *target* vertex v_t as follow:

$$\phi_{v_t, v_c} = \exp\left(\langle \omega_{v_c}, \omega_{v_t} \rangle \right)$$

$$\ell_{\text{random walk}} = \sum_{v_t, v_c} \log \left(\frac{\phi_{v_t, v_c}}{\sum_{k \in V} \phi_{v_k, v_t}} \right) \qquad (2)$$

In here, $\langle \cdot, \cdot \rangle$ denotes the inner product; ω_v denotes the real vector representation of node v; and W^{emb} is the output matrix containing all real vector representations. Although the log-likehood given by equation (1) is tractable, computing the normalization factor still remains an intractable task. Therefore, several approximation methods such as hierarchical softmax Morin and Bengio [2005]; Perozzi *et al.* [2014] or noise contrastive estimation Gutmann and Hyvärinen [2010]; Grover and Leskovec [2016] were used to further boost the computational speed of the algorithms of this category. The obtained embedding matrix will be used as a feature matrix for various machine learning algorithms such as node classification (Multi-class Linear Regression model) or link prediction.

2.2 Semi-supervised Network Labeling

Planetoid, proposed by Yang *et al.* [2016], works slightly different to other Skipgram-based models. Instead of generating graph context only from the network structure, *Planetoid* also samples nodes based on training labels. Furthermore, *Planetoid* injects the network node's feature vectors for better embedding and node labeling results. *Planetoid* can be considered an improvement of *SemiEmb*, proposed in Weston *et al.*

[2012]. Another semi-supervised learning model similar to *Planetoid* is Graph Convolutional Networks (GCN) Kipf and Welling [2016]. GCN uses the graph convolutional operation as a transformation for feature vectors on a network. By stacking these convolutional operation into a neural network, the authors of GCN has been able to achieve remarkable accuracy in node classification and link prediction results compared with the previous unsupervised and semi-supervised algorithms. Moreover, the running time for GCN was shorter compared with other algorithms such as *Deepwalk* or *Planetoid*. In the following paragraphs, we present the details for GCN.

The convolution on a graph G of a function of the graph Laplacian g_θ (also a filter or a kernel) and a signal x is defined as:

$$g_\theta * x = g_\theta(\mathcal{L})x,$$

where the normalized Laplacian $\mathcal{L} = U\Lambda U^\top$; U is the Fourier basis and Λ is the frequencies of the graph. Graph convolution has been shown effective in processing graph-structured data, and also argued to be the generalization of convolutional networks Shuman *et al.* [2013]; Defferrard *et al.* [2016]; Kipf and Welling [2016]. In practice, given a graph where each node has a feature vector, we can treat the feature vector of the graph as signals. The output y of these "signals" filtered by g_θ on the graph is given by the graph convolution and its inverse:

$$y = g_\theta(U\Lambda U^\top)x = U(g_\theta(\Lambda)U^\top)x \qquad (3)$$

Computing equation (3) is computationally expensive due to the matrix multiplication and eigenvector decomposition operations. Therefore, fast estimation methods such as Chebyshev polynomial was suggested in Hammond *et al.* [2011]. Under the assumption $\lambda_{\max} \approx 2$, Kipf and Welling further proposed the linear approximation for filtering signals X on a graph:

$$y \approx \theta \left(I + D^{-1/2}AD^{-1/2} \right) \qquad (4)$$

Following the graph convolution approximation above, the GCN neural network model can be described by the forward computation and the loss function:

$$Z_{\text{forward}} = \text{softmax}\left(\hat{A}\,\text{ReLU}(\hat{A}XW^{(0)})W^{(1)} \right)$$
$$loss = - \sum_{l \in \mathcal{Y}_L} \sum_{f=1}^{F} Y_{lf}\ln Z_{lf} \qquad (5)$$

where W is the weight matrix for hidden layer [0] and output layer [1]; $\hat{A} = \tilde{D}^{-1/2}\tilde{A}\tilde{D}^{-1/2}$ is called the *Renormalization Trick* by the authors Kipf and Welling [2016]. The loss function here is the cross-entropy error over all labeled samples similar to the model proposed in Yang *et al.* [2016]. The neural network is trained using Adam Kingma and Ba [2014] with back-propagation.

2.3 Motif Conductance

As an generalization of *network conductance* (also graph conductance or Cheeger constant of a graph), Benson *et al.* proposed the concept of *motif conductance* in Benson *et al.*

Algorithm 1: Motif co-occurrence matrix generation

Data: Graph $G = (V, E)$
Input: isBinary, $\mathbf{m}$, $\mathcal{A}$
Output: $M^{\mathbf{m}}$
begin
  diam $\leftarrow$ Diameter($\mathbf{m}$, $\mathcal{A}$);
  $V \leftarrow G$.nodes();
  for $i \in nodes$ **do**
    $G' \leftarrow$ induced graph from BFS(node, diam);
    for $j \in G'$.nodes() **do**
      if *(i,j) satisfies* $\mathcal{A}$ **then**
        $c \leftarrow$ count $\mathbf{m}$ in G';
        if *isBinary* **then**
          $M_{i,j}^{\mathbf{m}} \leftarrow 1$ if $c > 0$ else 0;
        else
          $M_{i,j}^{\mathbf{m}} \leftarrow c$;
      else
        $M_{i,j}^{\mathbf{m}} \leftarrow 0$;
  return $M^{\mathbf{m}}$

[2016]. Similar to network conductance, motif conductance is a *score* for a cut $(S, \bar{S})$ targeting a motif $\mathbf{m}$:

$$\phi_{\mathbf{m}}(S) = \frac{\text{cut}_{\mathbf{m}}(S, \bar{S})}{\min[\text{vol}_{\mathbf{m}}(S), \text{vol}_{\mathbf{m}}(\bar{S})]},$$

where S is a node set in a network G; $\bar{S}$ is the complement of S; $\text{vol}_{\mathbf{m}}(S)$ is the number of motif instances that resides in S. Intuitively, minimizing the motif conductance is equivalent to minimizing the number of motifs $\mathbf{m}$ split by the cut. A motif is split when there is at least one anchor node in S and at least one in $\bar{S}$. Benson *et al.* then perform motif analysis and graph clustering based on this definition of motif conductance. Their result further confirms the structural role of motifs in a complex network. It also hinted that there is a strong motif structure within a community, which we can use as prior knowledge for solving problems on networks.

Definition 2.1. *Motif co-occurrence matrix*: Given a graph $G = (V, E)$, in which $v \in V$. The motif co-occurrence matrix of a motif $\mathbf{m}$ is given by:

$$M_{i,j}^{\mathbf{m}} = \sum_{(v, \chi_{\mathcal{A}}(v)) \in \mathbf{m}} \mathbf{1}(i, j \subset \chi_{\mathcal{A}}(v))$$

In here, $\mathcal{A}$ represents the anchor set; $(v, \chi_{\mathcal{A}}(v))$ represents pairs of node $v \in V_G$ and the other anchor nodes generated by $\chi_{\mathcal{A}}$. If the anchor node set $\mathcal{A}$ is empty, all motif co-occurrence is counted toward the motif co-occurrence matrix M. Otherwise, only nodes in the anchor set will be counted. Figure 3 illustrates the bi-fan motif and its anchor set. Algorithm 1 provides the detail for constructing a motif co-occurrence matrix.

3 Our Methods

In this section, we present the detail of our methods. Firstly, we propose the basis for the network motif selection from a network. Secondly, we present two approaches employing
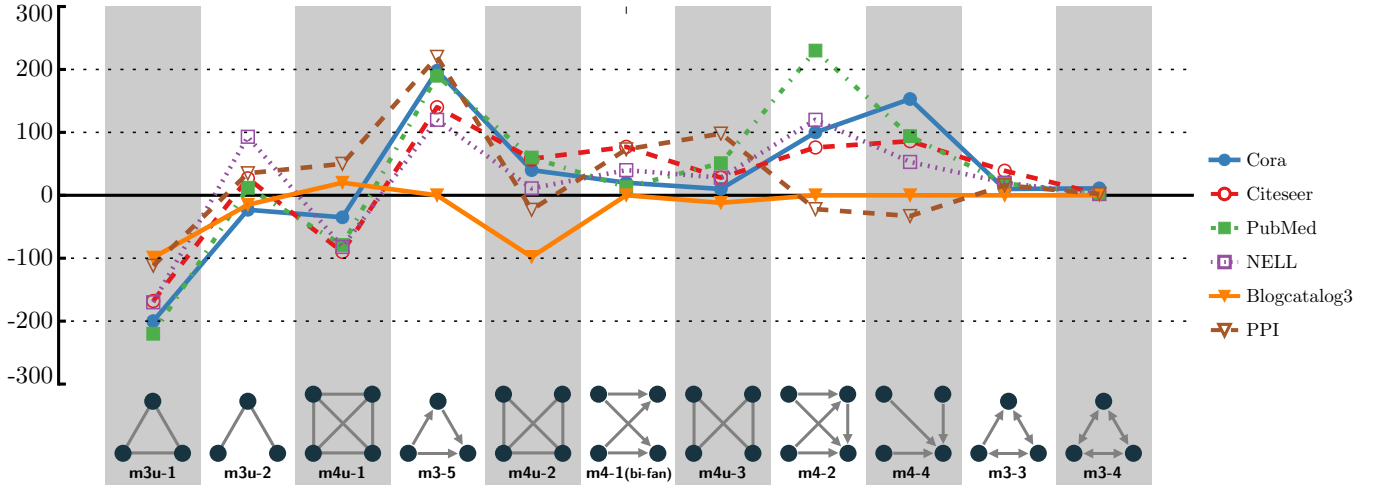
Figure 4: Significant graph for selected motifs (size-3 and size-4)

motif patterns to learn graph embeddings: *motifwalk* and *m-gcn*.

3.1 Motif Analysis and Motif Convolution

In the previous section, we have explained the importance of network motifs in network analysis. In this section, we present the metric for measuring network motif significance and the definition of motif Laplacian.

In order to measure the importance of a network motif, we compare the given network against a null model. The null model of an empirical network is an ensemble of randomly generated networks having the same number of nodes and edges as the network. For small networks with less than 10,000 edges, we generated 100 random networks as the ensemble of the null model. On the other hand, we generated 10 random networks for the null model of larger networks. The z-score is given by:

$$z\text{-score} = \frac{N_{\mathbf{m}}(G) - N_{\mathbf{m}}(G_{\text{random}})}{\sigma_{\mathbf{m}}(G_{\text{random}})}$$

where $N_{\mathbf{m}}(G)$ is the count of motif $\mathbf{m}$ in the empirical network; $N_{\mathbf{m}}(G_{\text{random}})$ is the mean of the null model; and $\sigma_{\mathbf{m}}(G_{\text{random}})$ is the variance. The z-score's values can range from $-\infty$ to $+\infty$. In practice, the most simple motifs (Figure 2 m3u-2, m3u-8 , m3u9) often have the highest frequencies and negative z-score. We ignored such motifs in our analysis. We select the motif which has the highest positive z-score because these motifs highlight the difference between the empirical network and random networks.

Convolution operations on a network can be viewed as a method to incorporate the nodes' information (e.g. feature vectors) and the network structure. The Fourier basis and network frequencies of a motif co-occurrence matrix is obtained through the eigenvalue decomposition of the motif Laplacian matrix $\mathcal{L}_{\mathbf{m}}$:

$$\mathcal{L}_{\mathbf{m}} = U_{\mathbf{m}} \Lambda_{\mathbf{m}} U_{\mathbf{m}}^{\top} \tag{6}$$

In here, $\Lambda_{\mathbf{m}} = diag(\lambda_{\mathbf{m}})$ is the frequencies of the motif net-

work; $\mathcal{L}_{\mathbf{m}}$ is the normalized motif Laplacian given by:

$$\mathcal{L}_{\mathbf{m}} = D_{\mathbf{m}}^{-1/2} L_{\mathbf{m}} D_{\mathbf{m}}^{-1/2}$$
$$= I - D_{\mathbf{m}}^{-1/2} M^{\mathbf{m}} D_{\mathbf{m}}^{-1/2} \tag{7}$$

where $L_{\mathbf{m}} = D_{\mathbf{m}} - M^{\mathbf{m}}$; $M^{\mathbf{m}}$ is the motif co-occurrence matrix; and $D_{\mathbf{m}} = \text{diag}(\sum_j M_{i,j}^{\mathbf{m}})$.

$U_{\mathbf{m}}$ from equation (7) is the Fourier basis of the network motif structure which is used for motif convolution. Given signal $x \in \mathbb{R}^n$ on a network, the motif Fourier transform is defined as $\hat{x} = U_{\mathbf{m}}^{\top} x$, and its inverse as $x = U_{\mathbf{m}} \hat{x}$. It follows that the output of signal x filtered by a function of graph Laplacian $g_{\theta}(L)$ parameterized by a set of coefficient θ is given by:

$$y = g_{\theta} * x = g_{\theta}(U_{\mathbf{m}} \Lambda_{\mathbf{m}} U_{\mathbf{m}}^{\top}) x$$
$$= U_{\mathbf{m}}(g_{\theta}(\Lambda_{\mathbf{m}}) U_{\mathbf{m}}^{\top}) x \tag{8}$$

Due to the complexity of eigenvector decomposition and matrix multiplication, we use the linear formulation suggested in Kipf and Welling [2016] to estimate the costly convolution operation:

$$y = g_{\theta'} * x \approx \sum_{k=0}^{K} \theta'_k T_k(\tilde{L}) x$$
$$\approx \theta'_0 x + \theta_1(L - I) = \theta'_0 x - \theta'_1 D_{\mathbf{m}}^{-1/2} M^{\mathbf{m}} D_{\mathbf{m}}^{-1/2} \tag{9}$$
$$\approx \theta \left(I + D_{\mathbf{m}}^{-1/2} M^{\mathbf{m}} D_{\mathbf{m}}^{-1/2} \right)$$

3.2 Biased Random Walk

Previous Skipgram-based graph embedding models employ random walks for graph context generation. To improve the embedding results, structure-aware context generation methods were proposed in Tang *et al.* [2015]; Grover and Leskovec [2016]. However, the limitation of *LINE* lies at the fact that it only consider the second-order proximity (bi-fan motif), and *node2vec* requires the costly cross-validation grid search for its hyper-parameters p and q. To solve the above-mentioned

problems, we propose a biased random walk algorithm for graph context generation which can be considered the generalization of *LINE* and *Deepwalk*. Since our algorithm decides the walking pattern supported by the most significant network motif before performing context generation, our method has the simplicity of *Deepwalk* while having the structure-aware context as of *LINE* and *node2vec*.

Our *motifwalk* algorithm has two steps: motif adjacency matrix construction and context generation. Firstly, we construct a binary motif co-occurrence matrix from the given network. We select the motif pattern as described in the previous section. Since the constructed matrix accounts the co-occurrence of network node pairs in a motif, it is a symmetric matrix. Secondly, after having the motif adjacency matrix, we run random walks on this new network induced by the adjacency matrix for motif context generation. The obtained motif context is used jointly with random walks context generated with the original network to train an embedding Skipgram model. Algorithm 1 and algorithm 2 describe the *motifwalk* algorithm.

Algorithm 2: Motif-aware graph context generation

Data: Graph $G = (V, E)$, Motif Graph $G_{\mathbf{m}} = (V, E)$
Input: length, nwalk, nmwalk
Output: context
begin
    context ← [ ];
    V ← G.nodes();
    nodes ← Shuffle(V);
    for *node $\in$ nodes* **do**
        walks ← [ ];
        for *i=0; i < nwalk; ++i* **do**
            walks += RandomWalk(graph=G,
                start=node, len=length);
        for *i=0; i < nmwalk; ++i* **do**
            walks += RandomWalk(graph=$G_{\mathbf{m}}$,
                start=node, len=length);
        context += walks;
    return *context*

Similar to other Skipgram-based models, *motifwalk* is an unsupervised algorithm which learns graph embedding through an optimization process. Since there is two network contexts generated in our algorithm, the objective function is given by:

$$\mathcal{O} = \underset{W^{\text{emb}}}{\arg\max} \left(\gamma \ell_{\text{random walk}} + (1 - \gamma) \ell_{\text{motif walk}} \right) \quad (10)$$

where $\ell_{\text{random walk}}$ and $\ell_{\text{motif walk}}$ are the log-likelihoods of the network contexts generated by random walk and motif walk respectively; γ is a hyper-parameter controlling the ratio between random walk. γ is selected empirically based on the ratio between edges in the motif network and edges in the given network. The likelihood of a *target* vertex v_t in the context of vertex v_c is given in equation (2). We use negative sampling with noise contrastive estimation loss as suggested in Mikolov and Dean [2013] for normalization factor estimation. The output of *motifwalk* is a set of real vectors representing each node in the given network. These vectors encode the underlying structural relationship between network nodes and can be used as feature vectors for link prediction and node classification.

3.3 Motif Convolutional Architecture

In the previous section, we have defined the graph convolution operation on motif co-occurrence matrix. We use the motif convolution as the second layer in our motif convolutional network (m-gcn). Based on the linear approximation proposed by Kipf and Welling, we define the forward computation of our model as:

$$
\begin{aligned}
Z_{\text{forward}} &= f(X, A, M) \\
&= \text{softmax}(\hat{M}\text{ReLU}(\hat{A}XW^{(0)})W^{(1)}),
\end{aligned} \quad (11)
$$

where A and M is a binary adjacency matrix and motif co-occurrence matrix respectively; $\hat{A}$ and $\hat{M}$ are constructed by the *renormalization trick* as suggested in Kipf and Welling [2016]; X contains the feature vectors for each graph node; $W^{(0)}$ and $W^{(1)}$ are learnable variables. We use Adam Kingma and Ba [2014] with back-propagation to optimize the cross-entropy loss E:

$$E = -\sum_{l \in \mathcal{Y}_L} \sum_{f=1}^{F} Y_{lf} \ln Z_{lf} \quad (12)$$

This formulation is the same as equation (5). The weight W is updated by the following rule:

$$\frac{\partial E}{\partial W_{i,j}^{(k)}} = \sum_{s=1}^{S} [x]^{\top} \frac{\partial E}{\partial y_{k,j}}, \quad (13)$$

where $W^{(k)}$ is the weight matrix of layer k; $y_{k,j}$ is the j^{th} output feature map of the sample s.

4 Experiments

To compare *motifwalk* performance with other unsupervised graph embedding models, we have chosen the node labeling task on various type of networks (with ground-truth node labels). We compare the performance of *motifwalk* to *Spectral Clustering*, *Deepwalk* Perozzi *et al.* [2014], *node2vec* Grover and Leskovec [2016]. Table 1 gives the networks' statistics. The second model *m-gcn* is a semi-supervised model for node labeling task. We present the comparison of our model with *planetoid* and *gcn* under the same experimental setting in Kipf and Welling [2016]. Table 2 shows the details of each dataset.

Blogcatalog3 Tang and Liu [2009]: This network is a blogger social network. Each node in the network represents an user. The node labels represent a blogger's interests, each node has at least one label. Since Blogcatalog is a social network, the results obtained from motif analysis agrees with the term "friends of a friend are friends". [1]

---

[1]Although the motif z-score for blogcatalog low, we select the triangle as a mean of applying a common knowledge into a friendship network. Further study for this particular z-score case will be conducted in the future.

DATASET	#CLASSES	#NODES	#EDGES	TRAINING RATIO
BLOGCATALOG	39	10,312	333,983	0.5
CITESEER	6	3,327	4,732	0.5

Table 1: Datasets for unsupervised embeddings

DATASET	#CLASSES	#NODES	#EDGES	#FEATURES
CITESEER	6	3,327	4,732	3,703
CORA	7	2,708	5,429	1,433
PUBMED	3	19,717	44,338	76,584
NELL	210	65,755	266,144	5,414

Table 2: Datasets for semi-supervised embeddings

Citeseer, Cora, Pubmed Sen *et al.* [2008]: These are citation networks consist of scientific papers and their citations represented by directed edges. Each node in these networks represents a scientific paper, associated with a tf-idf feature vector describing the content of the paper. Each paper also has exactly one label.

NELL Carlson *et al.* [2010]: This is a knowledge network where entities are connected by directed, labeled edges (relations). We use the preprocessing procedure as suggested in Kipf and Welling [2016] and Yang *et al.* [2016].

To select the most significant motif for each network, we compute the z-scores for size-3 and size-4 directed motifs and report the motif significance in Figure 4. The motif used for each graph is given in Table 3.

5 Results

Motifwalk algorithm is performed with following parameters: Number of walks per node ($nwalk$): 8; Walk length ($length$): 80; Skip-gram window size: 10; Samples per target: 4; Number of motif walks per node ($nmwalk$): 2. The learned embeddings are used to train a Multi-class Linear Regression model with train-test split ratio of 0.5 (for all dataset in table 1). We report the f1-macro score for comparison.

m-GCN model consists of two convolutional layers as described by equation (11). We train the model using Adam Kingma and Ba [2014] optimizer and dropout regularization for each layer. The accuracy scores for single label classification are given in Table 5.

6 Discussion

Our results in Table 5 have shown the effectiveness of using network motifs in learning network embeddings and classifying network nodes. The connection between two of our

DATASET	MOTIF
BLOGCATALOG	Figure 2 m3u-1
CITESEER	Figure 3 m4-2
CORA	Figure 3 m4-2
PUBMED	Figure 3 m4-2
NELL	Figure 2 m3-5

Table 3: Target motifs for each network

METHOD	BLOGCATALOG	CORA
Spectral clustering	0.23	0.67
Deepwalk	0.22	0.66
Node2Vec	0.22	0.66
motifwalk	**0.24**	**0.68**

Table 4: F1-macro score for multi-class labeling

METHOD	CITESEER	CORA	PUBMED	NELL
Deepwalk	43.2	67.2	65.3	58.1
motifwalk	45.7	68.0	64.9	58.8
Planetoid	64.7	75.7	77.2	61.9
GCN	70.3	81.5	79.0	66.0
m-GCN	**71.2**	**82.1**	**79.5**	**66.1**
m-GCN (rand. splits)	70.2 ± 0.5	81.1 ± 0.5	79.3 ± 0.7	62.0 ± 1.4

Table 5: Accuracy score for multi-class labeling

approaches is that both of them utilizes the network co-occurrence matrix. *Motifwalk* performs random walks on the induced motif network to generate a network context that is biased toward the targeted motif. This graph context generation algorithm can be seen as the generalization of *Deepwalk* and *LINE*. On the other hand, *m-gcn* uses the motif network as a Fourier basis for graph convolution. As a result, both approaches yields improvements compared with state-of-the-art algorithms on benchmark datasets. Furthermore, based on the result here, we believe the *m-gcn* model can also be used as mean to measure a motif's significant similar to motif conductance metric proposed by Benson *et al.*.

The weakness of our algorithms lies at building the motif co-occurrence matrix. Algorithms involving network motifs have high time complexity due to the problem of graph isomorphism Tran *et al.* [2014]. For such reason, in most large graph analysis, only motifs of size 5 or smaller are considered. In our experiments, we only consider motif of size 4 at most. This limitation is due to the large size of networks that we experimented. Although the analysis is limited by the motif size, we have been able to empirically show the effectiveness of the motif-aware methods. Furthermore, as mentioned in Benson *et al.* [2016], motif algorithms can be easily parallelized. Therefore, the extension to larger size motifs can be made possible by parallelizing the motif analysis procedures.

References

Uri Alon. Network motifs: theory and experimental approaches. *Nat. Rev. Genetics*, 8(6):450–461, 2007.

Gary D Bader, Doron Betel, and Christopher WV Hogue. Bind: the biomolecular interaction network database. *Nucleic acids research*, 31(1):248–250, 2003.

Albert-László Barabási. Network science book. *Network Science*, 2014.

Mikhail Belkin and Partha Niyogi. Laplacian eigenmaps and spectral techniques for embedding and clustering. In *NIPS*, volume 14, pages 585–591, 2001.

Austin R Benson, David F Gleich, and Jure Leskovec.

Higher-order organization of complex networks. *Science*, 353(6295):163–166, 2016.

Duncan S Callaway, Mark EJ Newman, Steven H Strogatz, and Duncan J Watts. Network robustness and fragility: Percolation on random graphs. *Phys. Rev. Let.*, 85(25):5468, 2000.

Shaosheng Cao, Wei Lu, and Qiongkai Xu. Grarep: Learning graph representations with global structural information. In *Procs. of the 24th ACM Int. Conf. on Information and Knowledge Management*, pages 891–900. ACM, 2015.

Andrew Carlson, Justin Betteridge, Bryan Kisiel, Burr Settles, Estevam R Hruschka Jr, and Tom M Mitchell. Toward an architecture for never-ending language learning. In *AAAI*, volume 5, page 3, 2010.

Michaël Defferrard, Xavier Bresson, and Pierre Vandergheynst. Convolutional neural networks on graphs with fast localized spectral filtering. In *Advances in Neural Information Processing Systems*, pages 3837–3845, 2016.

Aditya Grover and Jure Leskovec. node2vec: Scalable feature learning for networks. In *Proceedings of the 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, 2016.

Michael Gutmann and Aapo Hyvärinen. Noise-contrastive estimation: A new estimation principle for unnormalized statistical models. In *AISTATS*, volume 1, page 6, 2010.

David K Hammond, Pierre Vandergheynst, and Rémi Gribonval. Wavelets on graphs via spectral graph theory. *Applied and Computational Harmonic Analysis*, 30(2):129–150, 2011.

Christopher J Honey, Rolf Kötter, Michael Breakspear, and Olaf Sporns. Network structure of cerebral cortex shapes functional connectivity on multiple time scales. *Procs. of the Nat. Acad. of Sci.*, 104(24):10240–10245, 2007.

Marcus Kaiser and Claus C Hilgetag. Nonoptimal component placement, but short processing paths, due to long-distance projections in neural systems. *PLoS Comput Biol*, 2(7):e95, 2006.

Diederik Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.

Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *arXiv:1609.02907*, 2016.

Shmoolik Mangan and Uri Alon. Structure and function of the feed-forward loop network motif. *Procs. of the Nat. Acad. of Sci.*, 100(21):11980–11985, 2003.

T Mikolov and J Dean. Distributed representations of words and phrases and their compositionality. *Advances in neural information processing systems*, 2013.

Ron Milo, Shai Shen-Orr, Shalev Itzkovitz, Nadav Kashtan, Dmitri Chklovskii, and Uri Alon. Network motifs: simple building blocks of complex networks. *Science*, 298(5594):824–827, 2002.

F. Morin and Y. Bengio. Hierarchical probabilistic neural network language model. In *Aistats*, volume 5, pages 246–252. Citeseer, 2005.

Mark Newman. *Networks: an introduction*. Oxford university press, 2010.

Bryan Perozzi, Rami Al-Rfou, and Steven Skiena. Deepwalk: Online learning of social representations. In *Procs. of the 20th ACM SIGKDD Int. Conf. on Knowledge discovery and data mining*, pages 701–710. ACM, 2014.

Prithviraj Sen, Galileo Namata, Mustafa Bilgic, Lise Getoor, Brian Galligher, and Tina Eliassi-Rad. Collective classification in network data. *AI magazine*, 29(3):93, 2008.

David I Shuman, Sunil K Narang, Pascal Frossard, Antonio Ortega, and Pierre Vandergheynst. The emerging field of signal processing on graphs: Extending high-dimensional data analysis to networks and other irregular domains. *IEEE Signal Processing Magazine*, 30(3):83–98, 2013.

Lei Tang and Huan Liu. Relational learning via latent social dimensions. In *Proceedings of the 15th ACM SIGKDD Int. Conf. on Knowledge discovery and data mining*, pages 817–826. ACM, 2009.

Jian Tang, Meng Qu, Mingzhe Wang, Ming Zhang, Jun Yan, and Qiaozhu Mei. Line: Large-scale information network embedding. In *Procs. of the 24th Int. Conf. on WWW*, pages 1067–1077, 2015.

Ngoc Tam L Tran, Sominder Mohan, Zhuoqing Xu, and Chun-Hsi Huang. Current innovations and future challenges of network motif detection. *Briefings in bioinformatics*, page bbu021, 2014.

Martijn P Van Den Heuvel and Hilleke E Hulshoff Pol. Exploring the brain network: a review on resting-state fmri functional connectivity. *European Neuropsychopharmacology*, 20(8):519–534, 2010.

Robert West, Hristo S Paskov, Jure Leskovec, and Christopher Potts. Exploiting social network structure for person-to-person sentiment analysis. *arXiv preprint arXiv:1409.2450*, 2014.

Jason Weston, Frédéric Ratle, Hossein Mobahi, and Ronan Collobert. Deep learning via semi-supervised embedding. In *Neural Networks: Tricks of the Trade*, pages 639–655. Springer, 2012.

Zhilin Yang, William W. Cohen, and Ruslan Salakhutdinov. Revisiting semi-supervised learning with graph embeddings. In *Proc. of the 33rd Int. Conf. on Machine Learning*. ICML, 2016.